

COL380- Assignment 0


Name : Ujjwal Mehta | Entry No. : 2020CS10401

In this assignment we have learnt to use the profiling tool to profile and optimize our code and below is the corresponding description of each part.


2.1 Perf Stat

On running the given program by varying the number of threads used, we get the following plots.

1. The graph of the number of cycles with varying threads is as follows:



2. The graph of time elapsed with varying threads is as follows:



Now here above we observe that the peak for the number of cycles and the bottom for the time elapsed occurs when the number of threads are 16 and this happens because the number of CPUs inside the css cluster computers (css2 is the one on which I ran my program) are 16 so all the CPU gets utilized parallelly(hence maximum cycles and minimum time) while running the program and after that there is no significant change in graph since threads scheduled keep on switching on the CPUs as result of which we get the above graph.

2.2 Perf Record

We use the perf record command in order to generate the perf.data file which contains the cycles event that can be read using the perf report command using which we can obtain the information regarding the number of cycles, assembly code and the percentage of cpu which is used when we run the given classify file.

The assembly instructions which take the most CPU time are given in below screenshots which are indicated by the **red colour** .

		push	%rbx
		mov	0x10(%rdi),%rbp
		mov	0x18(%rdi),%ebx
		→ callq	omp_get_thread_num@plt
		mov	(%r12),%r8
		cmp	(%r8),%eax
		↓ jae	b0
		mov	%eax,%edi
		mov	0x8(%r8),%r10
		mov	0x8(%r12),%rsi
		mov	%eax,%r9d
		shl	\$0x2,%rdi
		mov	%eax,%ecx
		↓ jmp	70
		xchg	%ax,%ax
0.36	40:	cmp	0x4(%r11,%rax,8),%edx
37.77		↓ jg	93
0.39		shl	\$0x6,%rax
0.01		add	%rbp,%rax
0.02	4e:	mov	%r13d,0x4(%r12)
0.36		mov	(%rax),%rdx
0.22		cmp	%r9d,0x8(%rax)
		↓ jbe	b9
		lea	(%rdx,%rdi,1),%rax
		add	%ebx,%ecx
2.43		mov	(%rax),%edx

0.03		mov	(%r12),%edx
		test	%eax,%eax
		↓ jle	a8
0.02		mov	(%rsi),%r11
		lea	-0x1(%rax),%r14d
		xor	%eax,%eax
0.02	8a:	mov	%eax,%r13d
9.71		cmp	(%r11,%rax,8),%edx
13.88		↑ jge	40
18.66	93:	lea	0x1(%rax),%r13
1.29		cmp	%rax,%r14
		↓ je	a8
		mov	%r13,%rax
14.67		↑ jmp	8a
		nop	
	a8:	mov	%rbp,%rax
		xor	%r13d,%r13d

0.07	60:	nop	
		add	\$0x8,%rdx
	64:	cmp	%eax,0x4(%rcx)
51.65		↓ jne	81
		mov	0x18(%rbx),%r9
0.06		mov	(%rcx),%rcx
		mov	%esi,%r10d
		add	\$0x1,%esi
		add	(%r11),%r10d
0.03		mov	0x8(%r9),%r9
		mov	%rcx,(%r9,%r10,8)
0.23	81:	mov	%rdx,%rcx
		cmp	%rdx,%r8
47.96		↑ jne	60
	80:	add	%r12d,%eax

Yes, we can map the assembly instructions to the corresponding part of source code by just compiling our files with an additional **-g flag**.

3 Hotspot Analysis

After adding the **-g flag** we can clearly see the source code and the assembly code and upon looking at the annotate part, we can see that the hotspot lines of code are as follows:

		bool within(int val) const { // Return if val is within this range
		return(lo <= val && val <= hi);
0.42	40:	cmp 0x4(%r11,%rax,8),%edx
38.26		↓ jg 93
0.35		shl \$0x6,%rax
0.02		add %rbp,%rax
		_Z8classifyR4DataRK6Rangesj._omp_fn.0():
	4e:	mov %r13d,0x4(%r12)
		// and store the interval id in value. D is changed.
		counts[v].increase(tid); // Found one key in interval v
0.43		mov (%rax),%rdx
		_ZN7Counter8increaseEj():
		assert(id < _numcount);
0.28		cmp %r9d,0x8(%rax)
		↓ jbe b9
		_counts[id]++;
0.02		lea (%rdx,%rdi,1),%rax
		_Z8classifyR4DataRK6Rangesj._omp_fn.0():
		for(int i=tid; i<D.ndata; i+=numt) { // Threads together share-loop throu
		add %ebx,%ecx

	↓ jle a8
	if(_ranges[r].within(val))
0.02	mov (%rsi),%r11
	lea -0x1(%rax),%r14d
	xor %eax,%eax
0.03	8a: mov %eax,%r13d
	_ZNK5Range6withinEi():
	return(lo <= val && val <= hi);
9.59	cmp (%r11,%rax,8),%edx
14.58	↑ jge 40
	_ZNK6Ranges5rangeEib():
	for(int r=0; r<_num; r++) // Look through all intervals
18.25	93: lea 0x1(%rax),%r13
1.26	cmp %rax,%r14
	↓ je a8
	mov %r13,%rax
14.05	↑ jmp 8a
	nop

The prospective problem which makes this code snippet the top hotspot is that whenever we are finding the corresponding range in which we should assign a given data point then we call the **within method** the total number of ranges times due to which we have to load the address of this method a lot of times to execute it.

Yes, the code can be optimized in order to improve the performance of this hotspot and in order to optimize it quite significantly, we can change our algorithm a little bit to find the range corresponding to a given data point with a better time complexity like $O(\log N)$ rather than linear scanning using Tree like data structures to store ranges.

4 Memory Profiling

In this part of assignment, we will optimize our program further in order to make it more cache friendly by removing the instances of false sharing.

The top 2 hotspots obtained after running perf mem report on the original code are:

2.34	assert(id < _numcount); cmp %r9d,0x8(%rax) → jbe 33f9 <classify(Data&, Ranges const&, unsigned int) [clone ._omp _counts[id]++; lea (%rdx,%rdi,1),%rax _Z8classifyR4DataRK6Rangesj._omp_fn.0(): for(int i=tid; i<D.ndata; i+=numt) { // Threads together share-loop throu add %ebx,%ecx _ZN7Counter8increaseEj(): 87.59 mov (%rax),%edx add \$0x1,%edx mov %edx,(%rax) _Z8classifyR4DataRK6Rangesj._omp_fn.0(): 0.06 mov %ecx,%eax cmp %ecx,(%r8) → jbe 33f0 <classify(Data&, Ranges const&, unsigned int) [clone ._omp int v = D.data[i].value = R.range(D.data[i].key); // For each data, find t cltq lea (%r10,%rax,8),%r12 _ZNK6Ranges5rangeEib(): if(strict) { for(int r=0; r<_num; r++) // Look through all intervals if(_ranges[r].strictlyin(val)) return r; } else { 0.02 for(int r=0; r<_num; r++) // Look through all intervals mov 0x8(%rsi),%eax _Z8classifyR4DataRK6Rangesj._omp_fn.0(): 6.92 mov (%r12),%edx
------	---

0.03	the rcount = 0; xor %esi,%esi D2.data[rangeount[r-1]+rcount++] = D.data[d]; // Copy it to the appropri lea -0x4(%r14,%rdx,4),%r11 lea 0x8(%rcx),%rdx lea (%rdx,%r12,1),%r8 → jmp 3304 <classify(Data&, Ranges const&, unsigned int) [clone ._omp nop add \$0x8,%rdx if(D.data[d].value == r) // If the data item is in this interval cmp %eax,0x4(%rcx) 97.28 → jne 3321 <classify(Data&, Ranges const&, unsigned int) [clone ._omp D2.data[rangeount[r-1]+rcount++] = D.data[d]; // Copy it to the appropri 0.32 mov 0x18(%rbx),%r9 mov (%rcx),%rcx mov %esi,%r10d add \$0x1,%esi 0.22 add (%r11),%r10d 1.14 mov 0x8(%r9),%r9 mov %rcx,(%r9,%r10,8) for(int d=0; d<D.ndata; d++) // For each interval, thread loops through a 0.75 mov %rdx,%rcx
------	--

Now based on the above hotspot obtained, we can clearly identify 2 issues in the code which are leading to instances of **false sharing** and these are the lines

```
counts[v].increase(tid); // Found one key in interval v
```

```
D2.data[rangeount[r-1]+rcount++] = D.data[d]; // Copy it to the  
appropriate place in D2.
```

And the above lines there is false sharing because the threads are accessing the elements of the array in an **interleaved manner** which is leading to writing in same line by 2 different threads resulting in a lot of cache misses.

These instances of false sharing can be handled by making each thread process a contiguous section of the array hence making them independent from each other, hence we create our own matrix for the counts rather than using the counter object since where each thread will operate on a row increasing the cache hit rate and similarly removing interleaving inside the lower for loop as well.

After performing the above optimizations, we get the following report screenshots:

0.07	<pre> nop int rcount = 0; for(int d=0; d<D.ndata; d++) // For each interval, thread loops through test %r10d,%r10d → je 3346 <classify(Data&, Ranges const&, unsigned int) [clone ._om mov 0x8(%r13),%rdx int rcount = 0; xor %esi,%esi lea 0x8(%rdx),%rax </pre>
	<pre> lea (%rax,%r11,1),%r8 </pre>
0.01	<pre> → jmp 331c <classify(Data&, Ranges const&, unsigned int) [clone ._om nop add \$0x8,%rax if(D.data[d].value == r) // If the data item is in this interval cmp %ecx,0x4(%rdx) 84.06 → jne 333b <classify(Data&, Ranges const&, unsigned int) [clone ._om D2.data[rangeount[r-1]+rcount++] = D.data[d]; // Copy it to the approp 1.51 mov 0x18(%rbx),%r14 mov (%rdx),%rdx mov %esi,%r15d add \$0x1,%esi 0.25 add -0x4(%rbp,%rcx,4),%r15d mov 0x8(%r14),%r14 mov %rdx,(%r14,%r15,8) 0.17 for(int d=0; d<D.ndata; d++) // For each interval, thread loops through mov %rax,%rdx cmp %rax,%r8 → jne 3318 <classify(Data&, Ranges const&, unsigned int) [clone ._om 0.05 mov 0x8(%rdi),%esi for(int r=tid*(R.num()/numt); r<(tid+1)*(R.num()/numt); r++) { // Thread mov %esi,%eax xor %edx,%edx </pre>

0.02	<pre> mov 0x8(%r13),%rdx int rcount = 0; xor %esi,%esi lea 0x8(%rdx),%rax lea (%rax,%r11,1),%r8 → jmp 339c <classify(Data&, Ranges const&, unsigned int) [clone ._om nop add \$0x8,%rax if(D.data[d].value == r) // If the data item is in this interval cmp %ecx,0x4(%rdx) 0.02 → jne 33bb <classify(Data&, Ranges const&, unsigned int) [clone ._om 13.56 D2.data[rangecount[r-1]+rcount++] = D.data[d]; // Copy it to the appropri 0.04 mov 0x18(%rbx),%r9 mov (%rdx),%rdx mov %esi,%r10d add \$0x1,%esi 0.09 add -0x4(%rbp,%rcx,4),%r10d mov 0x8(%r9),%r9 mov %rdx,(%r9,%r10,8) for(int d=0; d<D.ndata; d++) // For each interval, thread loops through </pre>
------	--

Hence we see a significant decrease in terms of usage percentage.

Finally after running perf record with cache-misses flag on original as well as optimized code, we see that there is an increase in our cache hit rate.